

1.1 What is Computational Learning Theory

Machine Learning from theory point of view. What is ML?

- Automatic extraction of useful information from "raw" data
- Programs which improve performance via interaction with data
- Examples
 - Classification/clustering: text categorization, fraud detection, web search
 - Prediction: weather, financial market, earthquakes
 - Writing complex software: driving a car, etc.

1.2 Overview of Computational Learning Theory

- Specifying learning models
- Proving results in these models
- A learning model should specify
 - Who is learning? → computer program, usually restricted:
 - * polynomial learning time; $O(n^2), O(n \log n)$
 - * small sample complexity, particular format for output hypothesis
 - What are we learning?
 - * Skill (e.g. curve ball)
 - * Environment (e.g. NYC)
 - * In this class, we will focus on classification rules/Boolean functions
 - How does learner get information
 - * Learner is given examples $(x, f(x))$
 - x : feature vector
 - $f(x)$: label
 - * Passive learning (main focus of this lecture)
 - x can be chosen randomly from some distribution, or
 - some data $(x, f(x))$ can be chosen maliciously (corrupted data), or
 - x can be chosen by a helpful teacher
 - * Active learning: learner can make guess

- Learner chooses x , get $f(x)$ (Membership query)
 - Learner poses hypothesis, get counterexamples
 - Others (subset query, etc)
- Is information/data every noisy/incomplete?
 - * A few bits of $(x, f(x))$
 - * (x, y) , where y might have some noise
- What prior info does learner have?
 - * Typical assumption: there is a prior representation scheme for the function being learned. (e.g. "target concept" is known/assumed to be an "And" function of features)
 - * Performance criteria (Batch/offline vs online)
 - * How do we measure accuracy of learner's hypothesis
- **The main focus in this class** is learning boolean functions (binary classification rules/concepts) "concept learning"

1.3 Topics in this class

1.3.1 Our models

- Online mistake-bound model
- Probably Approximately Correct (PAC) model - which is the standard model for statistical learning
- Query models (Statistical Query, Membership Query, Equivalence Query, etc.)

1.3.2 Topics in these models

- Describe and analyze specific learning algorithms
- General theory
 - Necessary, sufficient conditions for learning
 - Sample complexity: how much data is need to get a "pre-specified" accuracy (e.g. PAC learning: exact characterization)
- Computational issues in learning (e.g. computational hardness)
- Learning from noisy data
- Boosting
- Compare learning models (e.g. Statistical Query learning → PAC learning)

1.4 Basic terms and concepts

1.4.1 Basics

- X = “instance space” = domain of the functions we’re learning (e.g. $X = \{\text{all cars in the world}\}$)
- We want to learn an unknown “target function” (or concept)
 - $c : X \rightarrow \{0, 1\}$ (e.g. $c(x) = 1$ if x is a mid-sized car or 0 if not)
- C : a “concept class”, which is a set of concepts (=set of subset of X)
 - e.g. $C = \{c_1 = \{\text{red cars}\}, c_2 = \{\text{convertibles}\}, \dots\}$

1.4.2 Basic idea of our learning models

- X and C are known to the learner
- There is a unknown target concept $c \in C$.
- Learner has some source of information about c , wants to identify/approximate c .
- Typically $X = \{0, 1\}^n$ or $X = \mathbb{R}^n$
- An $x \in X$ is $x = (x_1, \dots, x_n)$, where each x_k are features
- Learner “knows” C and X , but does not know target concept $c \in C$
- Examples
 - 1. Monotone conjunctions
 - $X = \{0, 1\}^n$ Boolean hypercube
 - C is all monotone (no negation) conjunctions (AND) over x_1, \dots, x_n (e.g., $c(x) = x_3 \wedge x_5 \wedge x_6$, then if $n = 7$ and $x = 1100111$, $c(x) = 0$ (False))
 - $|C| = 2^n$
 - 2. Conjunctions
 - $X = \{0, 1\}^n$ Boolean hypercube
 - C is all conjunctions (AND) over x_1, \dots, x_n (e.g., $\bar{x}_2 \wedge \bar{x}_4 \wedge \bar{x}_7 \wedge \bar{x}_8$)
 - $|C| = 3^n$
 - 3. Disjunctive Normal Form
 - $X = \{0, 1\}^n$ Boolean hypercube
 - C = all DNF formulas with $\leq n^2$ terms
 - A DNF = an OR of ANDs (e.g. $c(x) = (x_2 \wedge x_3) \vee (x_1 \wedge \bar{x}_2 \wedge x_6) \vee (x_4 \wedge \bar{x}_5 \wedge x_7 \wedge \bar{x}_9)$)
 - 4. k -DNF
 - $X = \{0, 1\}^n$ Boolean hypercube

- $C = \text{all } k\text{-DNF formulas}$
 - * A k -DNF is a DNF in which each term (=AND of literals has $\leq k$ variables)
 - * e.g. $c(x) = x_1\bar{x}_2 \vee x_2\bar{x}_6 \vee x_3\bar{x}_5$ is a 2-DNF
 - A DNF = an OR of ANDs (e.g. $c(x) = (x_2 \wedge x_3) \vee (x_1 \wedge \bar{x}_2 \wedge x_6) \vee (x_4 \wedge \bar{x}_5 \wedge x_7 \wedge \bar{x}_9)$)
- 5. Linear Threshold Functions (LTf)
 - $X = \mathbb{R}^n$
 - $C = \text{all LTF over } \mathbb{R}^n$ (LTf is also called halfspace)
 - A function $c : \mathbb{R}^n \rightarrow \{0, 1\}$ is LTF if there are n real weights w_1, \dots, w_n and a threshold θ s.t. for all $x = (x_1, \dots, x_n) \in \mathbb{R}^n$

$$c(x) = \begin{cases} 1, & \text{if } w_1x_1 + \dots + w_nx_n \geq \theta \\ 0, & \text{otherwise} \end{cases}$$

- We also can study LTFs over $X = \{0, 1\}^n$, e.g.

$$MAJ(x) = \begin{cases} 1, & \text{if } \sum_{i=1}^n x_i \geq \frac{n}{2} \\ 0, & \text{otherwise} \end{cases}$$