

5.1 Perceptron

- Algorithm to learn linear threshold functions over \mathbb{R}^n .
- Let $\mathbb{1}[v^\top x \geq \theta]$ be the target halfspace, where $v = (v_1, \dots, v_n) \in \mathbb{R}^n$, $x = (x_1, \dots, x_n) \in \mathbb{R}^n$.
- *Assumption 1.* $\theta = 0$ (homogeneous / origin-centered)
Indeed, we can reduce any θ to $\theta = 0$ case by introducing a new variable $x_{n+1} = 1$ so that

$$v^\top x \geq \theta \iff v'^\top x' \geq 0$$

, where $v' = (v_1, \dots, v_n, -\theta)$, $x' = (x_1, \dots, x_n, x_{n+1})$

- *Assumption 2.* Each example $x \in \mathbb{R}^n$ that we get has $\|x\|_2 = 1$. Note that rescaling a point x by $\|x\|_2$ does not change which side of the hyperplane $w^\top x = 0$ on example is on.
- *Assumption 3.* We can assure $\|v\|_2 = 1$
- Perceptron maintains hypothesis vector w so that $h_w(x) = \mathbb{1}[w^\top x \geq 0]$

5.1.1 Perceptron Learning Algorithm

- Initially, $w = (0, \dots, 0)$.
- Update hypothesis as follows
 - If we get an example x right, no change.
 - If $w^\top x \geq 0$, but the true label is $v^\top x < 0$, update $w \leftarrow w - x$. (False positive)
 - If $w^\top x < 0$, but the true label is $+$, update $w \leftarrow w + x$. (False negative)
- Intuition, consider a false positive case, so we updated w as $w_{new} \leftarrow w_{old} - x$. Then, $w_{new}^\top x = (w_{old} - x)^\top x = w_{old}^\top x - 1 \leq w_{old}^\top x$, which shows that the update reduce the function value for the x .

5.1.2 Perceptron Convergence Theorem

Theorem 5.1.1 Suppose we run Perceptron to learn halfspace/LTF $v^\top x \geq 0$, where assumptions 1 3 hold. Let the margin $\delta = \min|v^\top x|$ over all examples x given to algorithm. Then, Perceptron makes $\leq \frac{1}{\delta^2}$ mistakes.

Lemma 5.1.2 After Perceptron algorithm has made M mistakes, we have $w^\top v \geq \delta M$.

Proof: We will show that each mistake increases $w^\top v$ by $\geq \delta$.

- Initially, $M = 0, w = (0, \dots, 0) \rightarrow$ the proposition trivially holds.
- If mistake is on a positive example x , $w_{new} = w_{old} + x$. Thus,

$$\begin{aligned} w_{new}^\top v &= (w_{old} + x)^\top v \\ &= w_{old}^\top v + x^\top v \\ &\geq w_{old}^\top v + \delta \quad \because \text{by the definition of } \delta \end{aligned}$$

Recall $\delta = \min|v^\top x|$ over all x in the input. Since x is positive, $v^\top x \geq 0$.

- If mistake is on a negative example x , $w_{new} = w_{old} - x$. Thus,

$$(w_{old} - x)^\top v = w_{old}^\top v - x^\top v \geq w_{old}^\top v + \delta$$

since $x^\top v < 0$ and $|x^\top v| \geq \delta$.

■

Lemma 5.1.3 *After M mistakes, $\|w\|_2 \leq M$.*

Proof:

- Initially, $M = 0$ and $\|w\|_2 = 0$, so the proposition trivially holds.
- If mistake is on positive example x ,

$$\|w + x\|_2^2 = \|w\|_2^2 + 2 \underbrace{w^\top x}_{<0} + \|x\|_2^2 < \|w\|_2^2 + 1$$

- If mistake is on negative example x ,

$$\|w - x\|_2^2 = \|w\|_2^2 - 2 \underbrace{w^\top x}_{>0} + \|x\|_2^2 < \|w\|_2^2 + 1$$

■

Proof of 5.2.1: Since $w^\top v = \|w\|_2 \underbrace{\|v\|_2}_{=1} \underbrace{\cos \theta(w, v)}_{\leq 1}$, $w^\top v \leq \|w\|_2$. Thus,

$$\delta M \leq w^\top v \leq \|w\|_2 \leq \sqrt{M}$$

From the inequality above, we can show $M \leq \frac{1}{\delta^2}$

■

Remarks

- Perceptron is
 - simple

- noise-tolerant
- kernelizable
- But in general $\frac{1}{\delta^2}$ bound is not always good. For example, for decision lists $\{0, 1\}^n$, $\delta = 2^{-\Theta(n)}$. There exist different algorithms (e.g. Ellipsoid algorithm) with mistake bound $\log(\frac{1}{\delta})^2$, but more complex, each update slower, not noise tolerant.
- Best of both worlds: Dunagan-Vempala's "Rescaled Perceptron"

5.2 Generic Algorithm and Bounds for Online Learning (Halving Algorithm)

Assume that the concept class C is finite. Halving algorithm learns any finite concept class C with $\leq \log_2 |C|$ mistakes.

5.2.1 Halving Algorithm Procedure

- Let $CONSIST \subseteq C$ by the set of all $c \in C$ which are consistent with all labeled examples seen so far.
- Initial $CONSIST = C$
- Given example x , the hypothesis that the algorithm uses is the majority vote over all concepts over $CONSIST$ on x .
- After you receive the true value $c(x)$, update $CONSIST$ so that only consistent hypotheses are included in $CONSIST$.

5.2.2 Analysis on Halving Algorithm

Theorem 5.2.1 *For any finite concept class C , Halving algorithm has the mistake bound $\leq \log_2 |C|$.*

Proof: Every mistake causes new value of $|CONSIST|$ to become at most half of old volume. ■

5.2.3 Examples of Halving Algorithm

Decision List

- $C = \{\text{all DLs over } x_1, \dots, x_n \text{ with length-}r\}$.
- Recall $|C| \leq (4n+2)^{r+1}$, which leads to $\log_2 |C| = O(r \log n)$ by halving algorithm. Note that out algorithm the lecture 2, 3 had the mistake bound $O(rn)$. It turns out that the upper bound by halving algorithm is optimal for any algorithm.
- Drawbacks
 - Very inefficient (time, space at least $|C|$).

- Not noise tolerant.
- In some cases, we can match this mistake upper bound with efficient algorithm. (e.g. disjunctions).

Delta function

- Halving algorithm can sometimes make even fewer mistakes than $\log_2 |C|$.
- Consider $X = \{1, \dots, N\}$,

$$\delta_i(j) = \begin{cases} 1, & \text{if } j = i \\ 0, & j \neq i \end{cases}$$

and $C = \{\text{all } N \text{ functions } \delta_i \text{ (singletons)}\}$. Then, halving algorithm makes 1 mistake.

Remarks

- In general, we don't have a computationally efficient analogue of halving algorithm.
- There exists noise tolerant version of halving algorithm.